

Threads synchronization

Contents

Thread basic

Thread mutex locking

Thread advance locking and condition variables

Atomic types

VC IVI Development Center Vietnam

Hanoi, November 2018

Threads

- ▶ Threads allow parallel execution of processes or distinct parts of a single process.

- ▶ All threads within a process share:
 - ◇ The same address space
 - ◇ Process instructions
 - ◇ Most data
 - ◇ Open files (descriptors)
 - ◇ Signals and signal handlers
 - ◇ Current working directory
 - ◇ User and group id

Threads

- ▶ Each thread has a unique:
 - ◇ Thread ID (tid)
 - ◇ Set of registers, stack pointer
 - ◇ Stack for local variables, return addresses
 - ◇ Signal mask
 - ◇ Priority
 - ◇ Return value: errno

- ▶ Basic thread operations:
 - ◇ Creation
 - ◇ Termination
 - ◇ Joining
 - ◇ Synchronization

Problems

In the general case, you often use shared objects between the threads. And when you do it, you will face another problem: synchronization.

Solution: provide functions that will block one thread if another thread is trying to access data that it is currently using.

- ▶ There are several ways to fix this problem:
- ▶ Semaphores
- ▶ Atomic references
- ▶ Monitors
- ▶ Condition codes
- ▶ Compare and swap
- ▶ etc.

~~Use a mutex to make a thread-safe~~

- ▶ There are two important methods on a mutex: `lock()` and `unlock()`. As their names indicate, the first one enable a thread to obtain the lock and the second releases the lock. The `lock()` method is blocking. The thread will only return from the `lock()` method when the lock has been obtained.

~~Automatic management of locks~~

It exists a good solution to avoid forgetting to release the lock: **std::lock_guard**.

This class is a simple smart manager for a lock. When the `std::lock_guard` is created, it automatically calls `lock()` on the mutex. When the guard gets destructed, it also releases the lock.

~~Recursive locking~~

The threads tries to acquire the lock again, but the lock is already locked. This is a case of deadlock. By default, a thread cannot acquire the same mutex twice.

Timed locking

Sometimes, you doesn't want a thread to wait ad infinitum for a mutex. For example, if your thread can do something else when waiting for the thread.

the standard library has a solution: **std::timed_mutex** and **std::recursive_timed_mutex**

You have access to the same functions as a **std::mutex**: *lock()* and *unlock()*, but you have also two new functions: *try_lock_for()* and *try_lock_until()*.

Condition variables

A condition variable manages a list of threads waiting until another thread notify them.

Each thread that wants to wait on the condition variable has to acquire a lock first. The lock is then released when the thread starts to wait on the condition and the lock is acquired again when the thread is awakened.

Atomic Types

The C++11 Concurrency Library introduces Atomic Types as a template class: `std::atomic`.

The main advantage of this technique is its performance.

`std::atomic` is specialized for all integral types to provide member functions specific to integral.

Thank you!